

The Machine Learning Pipeline (workflow)

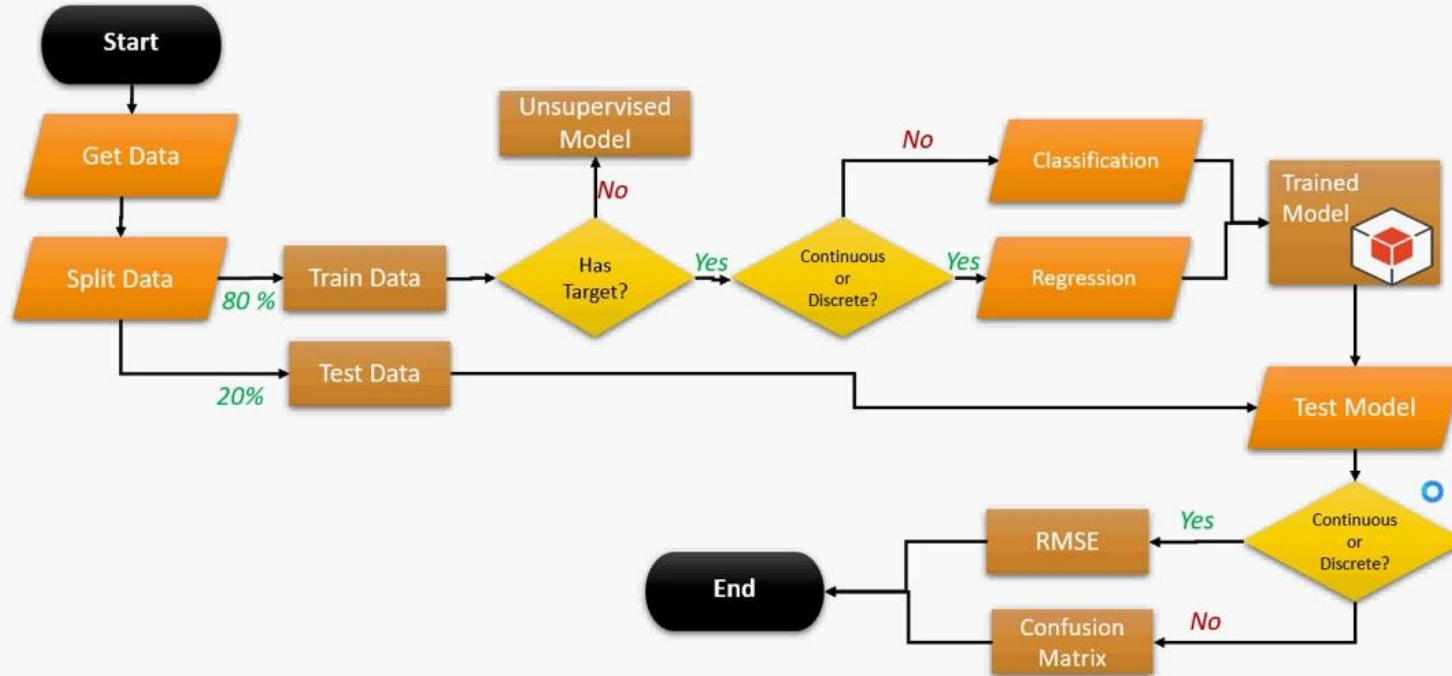
Marc Berghouse and Lazaro Perez

The Machine Learning Pipeline

- The pipeline is composed of all of the computation and analysis required to get from the beginning to the end of a machine learning project
- Gather data
- Clean and preprocess data
- Exploratory data analysis
- Split data into training, validation and testing sets (holdout vs cross-validation)
- Develop/import model
- Train model
- Test model
- Hyperparameter tuning
- Report and visualize results



Machine Learning – Model Flowchart



Gather Data

- We need data
- Good data is hard to find
 - No cure all method to easily finding and extracting data
 - Be persistent
- Data almost always needs to be cleaned and preprocessed
 - Think about, how can I make this data look like an Excel spreadsheet/Pandas dataframe
- Generally try to read in data with `pd.read_csv()`

Clean and Preprocess data

- Read data and parse into a pandas dataframe or numpy array
- Take care of NAN values and outliers
 - Imputation or dropping
 - NaN values propagate through calculations
 - Outlier detection (kNN, support vector machine, covariance)
- Use sklearn's `StandardScaler()` to scale your data to mean of 0 with a standard deviation of 1
- Make necessary data transformations
 - For dates, `pd.to_datetime()` should work, but I've had issues with it
 - You can also count number of days or hours from the beginning of the dataset and use this as the date/time variable
 - Units should be in the same system (kg, m, s; lb, ft, s)
 - Encode string-type variables
 - Transformations should be done before using `StandardScaler()`

Encoding

- Convert categorical variables (strings) to numbers

```
>>> enc = preprocessing.OrdinalEncoder(encoded_missing_value=-1)
>>> X = [['male'], ['female'], [np.nan], ['female']]
>>> enc.fit_transform(X)
array([[ 1.],
       [ 0.],
       [-1.],
       [ 0.]])
```

```
>>> import numpy as np
>>> from sklearn.preprocessing import FunctionTransformer
>>> transformer = FunctionTransformer(np.log1p, validate=True)
>>> X = np.array([[0, 1], [2, 3]])
>>> # Since FunctionTransformer is no-op during fit, we can cal
>>> transformer.transform(X)
array([[0.          , 0.69314718],
       [1.09861229, 1.38629436]])
```

Fill Missing or NaN Values (sklearn)

```
>>> import numpy as np
>>> from sklearn.impute import KNNImputer
>>> nan = np.nan
>>> X = [[1, 2, nan], [3, 4, 3], [nan, 6, 5], [8, 8, 7]]
>>> imputer = KNNImputer(n_neighbors=2, weights="uniform")
>>> imputer.fit_transform(X)
array([[1. , 2. , 4. ],
       [3. , 4. , 3. ],
       [5.5, 6. , 5. ],
       [8. , 8. , 7. ]])
```

sklearn.impute.IterativeImputer

```
class sklearn.impute.IterativeImputer(estimator=None, *, missing_values=nan, sample_posterior=False, max_iter=10,
tol=0.001, n_nearest_features=None, initial_strategy='mean', imputation_order='ascending', skip_complete=False, min_value=-
inf, max_value=inf, verbose=0, random_state=None, add_indicator=False, keep_empty_features=False) \[source\]
```

Multivariate imputer that estimates each feature from all the others.

```
>>> # explicitly require this experimental feature
>>> from sklearn.experimental import enable_iterative_imputer # noqa
>>> # now you can import normally from sklearn.impute
>>> from sklearn.impute import IterativeImputer
```

```
>>> import pandas as pd
>>> df = pd.DataFrame([["a", "x"],
...                    [np.nan, "y"],
...                    ["a", np.nan],
...                    ["b", "y"]], dtype="category")
>>> imp = SimpleImputer(strategy="most_frequent")
>>> print(imp.fit_transform(df))
[['a' 'x']
 ['a' 'y']
 ['a' 'y']
 ['b' 'y']]
```

6.4.7. Estimators that handle NaN values

Some estimators are designed to handle NaN values without prep type (cluster, regressor, classifier, transform):

- **Estimators that allow NaN values for type regressor:**
 - [HistGradientBoostingRegressor](#)
- **Estimators that allow NaN values for type classifier:**
 - [HistGradientBoostingClassifier](#)
- **Estimators that allow NaN values for type transformer:**
 - [IterativeImputer](#)
 - [KNNImputer](#)
 - [MaxAbsScaler](#)
 - [MinMaxScaler](#)
 - [MissingIndicator](#)
 - [PowerTransformer](#)
 - [QuantileTransformer](#)
 - [RobustScaler](#)
 - [SimpleImputer](#)
 - [StandardScaler](#)
 - [VarianceThreshold](#)

Fill Missing or NaN Values (Pandas)

```
In [54]: dff.fillna(dff.mean())
```

```
Out[54]:
```

	A	B	C
0	0.271860	-0.424972	0.567020
1	0.276232	-1.087401	-0.673690
2	0.113648	-1.478427	0.524988
3	-0.140857	0.577046	-1.715002
4	-0.140857	-0.401419	-1.157892
5	-1.344312	-0.401419	-0.293543
6	-0.109050	1.643563	-0.293543
7	0.357021	-0.674600	-0.293543
8	-0.968914	-1.294524	0.413738
9	0.276662	-0.472035	-0.013960

```
In [55]: dff.fillna(dff.mean()["B":"C"])
```

```
Out[55]:
```

	A	B	C
0	0.271860	-0.424972	0.567020
1	0.276232	-1.087401	-0.673690
2	0.113648	-1.478427	0.524988
3	NaN	0.577046	-1.715002
4	NaN	-0.401419	-1.157892
5	-1.344312	-0.401419	-0.293543
6	-0.109050	1.643563	-0.293543
7	0.357021	-0.674600	-0.293543
8	-0.968914	-1.294524	0.413738
9	0.276662	-0.472035	-0.013960

Fill gaps forward or backward

Using the same filling arguments as [reindexing](#),
backward:

```
In [45]: df
```

```
Out[45]:
```

	one	two	three
a	NaN	-0.282863	-1.509059
c	NaN	1.212112	-0.173215
e	0.119209	-1.044236	-0.861849
f	-2.104569	-0.494929	1.071804
h	NaN	-0.706771	-1.039575

```
In [46]: df.fillna(method="pad")
```

```
Out[46]:
```

	one	two	three
a	NaN	-0.282863	-1.509059
c	NaN	1.212112	-0.173215
e	0.119209	-1.044236	-0.861849
f	-2.104569	-0.494929	1.071804
h	-2.104569	-0.706771	-1.039575

Filling missing values: fillna

`fillna()` can "fill in" NA values with non-NA data in a

Replace NA with a scalar value

```
In [42]: df2
```

```
Out[42]:
```

	one	two	three	four	five	t
a	NaN	-0.282863	-1.509059	bar	True	
c	NaN	1.212112	-0.173215	bar	False	
e	0.119209	-1.044236	-0.861849	bar	True	20
f	-2.104569	-0.494929	1.071804	bar	False	20
h	NaN	-0.706771	-1.039575	bar	True	

```
In [43]: df2.fillna(0)
```

```
Out[43]:
```

	one	two	three	four	five	t
a	0.000000	-0.282863	-1.509059	bar	True	
c	0.000000	1.212112	-0.173215	bar	False	
e	0.119209	-1.044236	-0.861849	bar	True	2
f	-2.104569	-0.494929	1.071804	bar	False	2
h	0.000000	-0.706771	-1.039575	bar	True	

```
In [44]: df2["one"].fillna("missing")
```

```
Out[44]:
```


a	missing
c	missing
e	0.119209
f	-2.104569
h	missing

Exploratory Data Analysis

- Understand and communicate what is going on with your data at a general level
- Pandas profiler
 - Basic stats for whole dataframe and individual variables
 - Scatter plots and correlation heatmap of all variables
 - If you can't install for whatever reason, you can use matplotlib to generate the same plots
- Feature Analysis
 - Unsupervised methods
 - Principal Component Analysis (PCA)
 - Outlier detection
 - Supervised methods
 - Linear Discriminant Analysis (LDA)
 - Random Forest
 - PCA, LDA, and factor analysis all look for linear combinations of variables that best explain the data

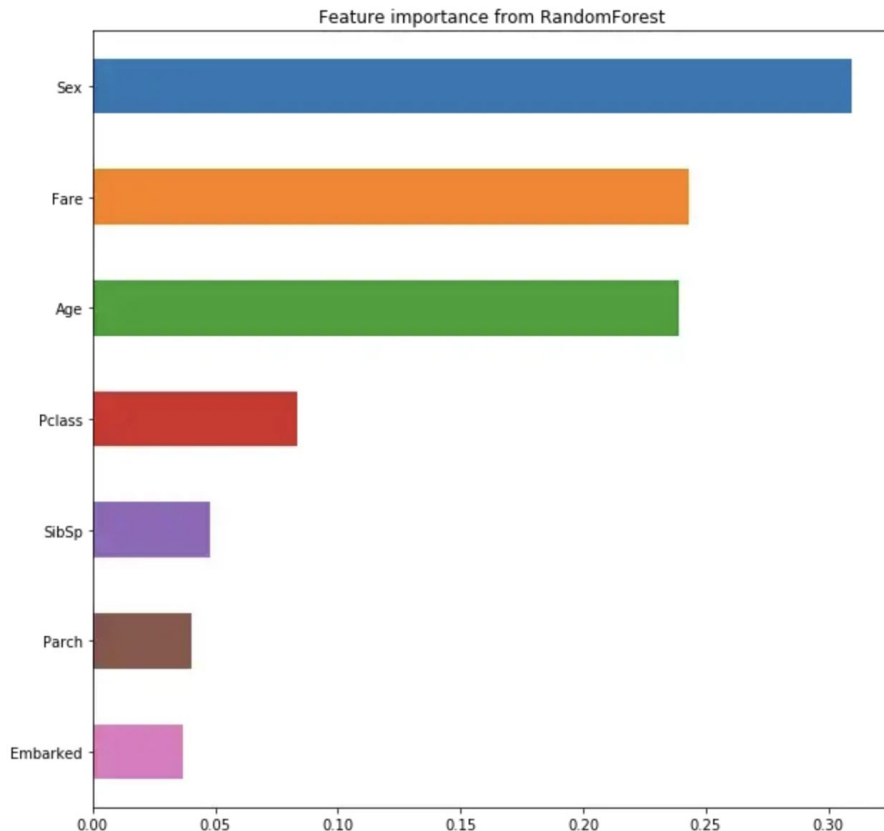
Principal Component Analysis (PCA)

- “Linear dimensionality reduction using Singular Value Decomposition of the data to project it to a lower dimensional space”
- Separates data into new dimensions based on the **combinations** of attributes that account for most of the variance in the data (explanatory power)
- **LDA** is like PCA but it looks for attributes that explain the most **variance between classes**



```
from sklearn.ensemble import RandomForestClassifier
rf_clf = RandomForestClassifier(n_estimators = 500, max_depth=12)
rf_clf.fit(X_train, y_train)
rf_y_pred = rf_clf.predict(X_val)
```

Random Forest Feature Engineering



Titanic Dataset

Name	Variable explanation
pclass	Passenger Class (1 = 1st;2 = 2nd;3 = 3rd)
Survived	Survival (0 = no, 1 = yes)
Name	Passenger name
Sex	Gender of passenger
Age	Age of passenger
Sibsp	(number of siblings/spouses aboard)
Parch	(number of parents/children aboard)
Ticket	Ticket number
Fare	Passenger fare (£)
Cabin	Cabin
Embarked	Port of Embarkation (C = Cherbourg; Q = Queenstown; S = Southampton)
Boat	Lifeboat
Body	Body Identification Number
Home.dest	Home/Destination

Training and testing splits

- We use training data to fit the model
- We use testing data (unseen during training) to test the model
 - Data leakage happens when the model somehow has access to testing samples during training
- Holdout
 - Only one train-test split
- Cross-Validation (CV)
 - Divide dataset into 6 parts
 - Train on five parts, test on 6th
 - Switch training/testing parts until all pairs are complete

```
>>> from sklearn.svm import SVC
>>> from sklearn.preprocessing import StandardScaler
>>> from sklearn.datasets import make_classification
>>> from sklearn.model_selection import train_test_split
>>> from sklearn.pipeline import Pipeline
>>> X, y = make_classification(random_state=0)
>>> X_train, X_test, y_train, y_test = train_test_split(X, y,
...                                                    random_state=0)
>>> pipe = Pipeline([('scaler', StandardScaler()), ('svc', SVC())])
>>> # The pipeline can be used as any other estimator
>>> # and avoids leaking the test set into the train set
>>> pipe.fit(X_train, y_train)
Pipeline(steps=[('scaler', StandardScaler()), ('svc', SVC())])
>>> pipe.score(X_test, y_test)
0.88
```



Model Development

- Choose the appropriate model to solve the problem
- Whole lecture on this later
- Classification or Regression?
- Supervised, unsupervised or reinforcement learning?
 - https://scikit-learn.org/stable/supervised_learning.html
 - https://scikit-learn.org/stable/unsupervised_learning.html
- Read the Sklearn documentation for the model you pick to make sure you are using it in the right situation and passing the correct arguments to the model function
- Hyperparameter tuning (more on this later)
 - https://scikit-learn.org/stable/modules/grid_search.html

Model Training

We use the `.fit()` method for all training

- We use **fit()** method only on the training data. Why? Because we don't know what our testing data (unseen data) is, hence using the **fit()** method on the test data would not give us a good estimate of how our model is performing.
- We use **transform()** method on train data as well as test data as we need to perform transformation in both cases.

In Case Of Transformers

Transformers are for pre-processing the data before modelling.

```
lin_reg.fit(X_train_diab, y_train_diab)
y_pred=lin_reg.predict(X_test_diab)
print ('R2: ',lin_reg.score(X_test_diab,y_test_diab))
print ('RMSE: ',mean_squared_error(y_pred,y_test_diab,squared=False))
```

```
R2: 0.4526066021617382
RMSE: 53.85325698491439
```

```
knn_class=skn.KNeighborsClassifier(n_neighbors=i)
knn_class.fit(X_train_iris, y_train_iris)
y_pred=knn_class.predict(X_test_iris)
```

- **fit ()** — This method goes through the training data, calculates the parameters (like mean (μ) and standard deviation (σ) in StandardScaler class) and saves them as internal objects.
- **transform()** — The parameters generated using the **fit()** method are now used and applied to the training data to update them.
- **fit _Transform()** — This method may be more convenient and efficient for modelling and transforming the training data simultaneously.

<https://medium.com/nerd-for-tech/difference-fit-transform-and-fit-transform-method-in-scikit-learn-b0a4efcab804>

Model Testing/Scoring

- We use the `.predict()` method to generate predictions of our testing data
- If we use `.predict()`, we then have to generate a score by comparing the predicted data (“`y_pred`”) with the true data (“`y_test`”)
- If we use `.score()`, Sklearn does the prediction and score generation together
- The `.score()` method uses the arguments (`x_test,y_test`), but if you call one of the specific scoring functions on the right, the arguments are (`y_pred,y_test`)

```
lin_reg.fit(X_train_diab, y_train_diab)
y_pred=lin_reg.predict(X_test_diab)
print ('R2: ',lin_reg.score(X_test_diab,y_test_diab))
print ('RMSE: ',mean_squared_error(y_pred,y_test_diab,squared=False))
```

R2: 0.4526066021617382

RMSE: 53.85325698491439

Scoring Classification	Function
'accuracy'	<code>metrics.accuracy_score</code>
'balanced_accuracy'	<code>metrics.balanced_accuracy_score</code>
'top_k_accuracy'	<code>metrics.top_k_accuracy_score</code>
'average_precision'	<code>metrics.average_precision_score</code>
'neg_brier_score'	<code>metrics.brier_score_loss</code>
'f1'	<code>metrics.f1_score</code>
'f1_micro'	<code>metrics.f1_score</code>
'f1_macro'	<code>metrics.f1_score</code>
'f1_weighted'	<code>metrics.f1_score</code>
'f1_samples'	<code>metrics.f1_score</code>
'neg_log_loss'	<code>metrics.log_loss</code>
'precision' etc.	<code>metrics.precision_score</code>
'recall' etc.	<code>metrics.recall_score</code>
'jaccard' etc.	<code>metrics.jaccard_score</code>
'roc_auc'	<code>metrics.roc_auc_score</code>



Regression	
'explained_variance'	<code>metrics.explained_variance_score</code>
'max_error'	<code>metrics.max_error</code>
'neg_mean_absolute_error'	<code>metrics.mean_absolute_error</code>
'neg_mean_squared_error'	<code>metrics.mean_squared_error</code>
'neg_root_mean_squared_error'	<code>metrics.mean_squared_error</code>
'neg_mean_squared_log_error'	<code>metrics.mean_squared_log_error</code>
'neg_median_absolute_error'	<code>metrics.median_absolute_error</code>
'r2'	<code>metrics.r2_score</code>
'neg_mean_poisson_deviance'	<code>metrics.mean_poisson_deviance</code>
'neg_mean_gamma_deviance'	<code>metrics.mean_gamma_deviance</code>
'neg_mean_absolute_percentage_error'	<code>metrics.mean_absolute_percentage_error</code>

Report and Visualize Results (Regression)

- Regression

- Scatter plots with line of best fit
- For multiple linear regression, you can use 3D plots (x,y, and z coordinates), color, and scatter point size to add extra dimensions to the plot
- Correlation matrix/heatmap (already done with Pandas profiler)
- Box or violin plots of scores for different models/experiments/runs
 - Test different models under a range of changing variables/parameters (range of scores for each model)
 - Test the range of results for probabilistic models (different result after each run, like with neural nets)

```
sns.boxplot(data=df, x="age", y="class")
```




Report and Visualize Results (Classification)

- Confusion Matrix
- ROC-AUC curve
 - False positive rate (FPR) vs true positive rate (TPR)


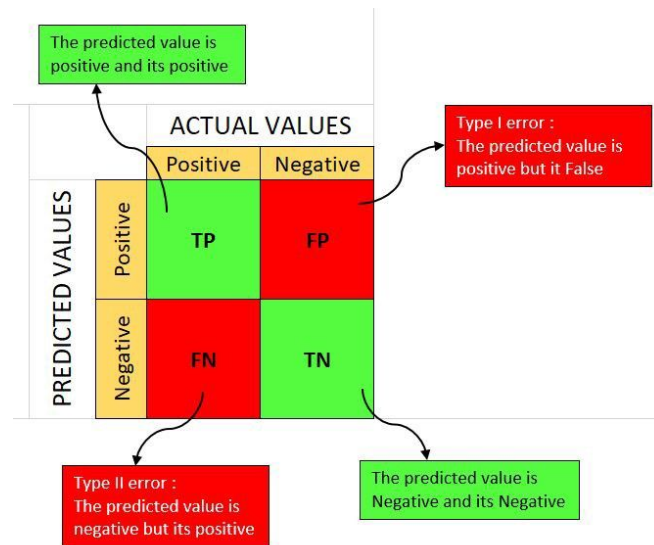
```
from sklearn.metrics import roc_curve
from sklearn.metrics import RocCurveDisplay
```

```
y_score = clf.decision_function(X_test)
```

```
fpr, tpr, _ = roc_curve(y_test, y_score, pos_label=clf.classes_[1])
roc_display = RocCurveDisplay(fpr=fpr, tpr=tpr).plot()
```



<https://medium.com/analytics-vidhya/what-is-a-confusion-matrix-d1c0f8feda5>



Example Walkthrough of Iris dataset confusion matrix and best nearest neighbors searcher

```
import sklearn.metrics as skm
best_i=0
equal_list=[]

for i in range(1,100):
    knn_class=skn.KNeighborsClassifier(n_neighbors=i)
    knn_class.fit(X_train_iris, y_train_iris)
    y_pred=knn_class.predict(X_test_iris)
    if i==1:
        best_mat=skm.confusion_matrix(y_pred,y_test_iris)
        best_score=(best_mat[0,0]+best_mat[1,1]+best_mat[2,2])/30
        best_i=1
    else:
        new_mat=skm.confusion_matrix(y_pred,y_test_iris)
        if (new_mat[0,0]+new_mat[1,1]+new_mat[2,2])/30>best_score:
            best_mat=skm.confusion_matrix(y_pred,y_test_iris)
            best_score=(new_mat[0,0]+new_mat[1,1]+new_mat[2,2])/30
            best_i=i
        elif new_mat[0,0]+new_mat[1,1]+new_mat[2,2]==best_score:
            equal_list.append(i)

print ('Classification Performance: ',best_score,' NN = ',str(best_i))
```

Classification Performance: 1.0 NN = 1

Help for Hw1 Q2

```
In [16]: from sklearn.datasets import load_iris, load_diabetes
iris = load_iris()
diab = load_diabetes()
from sklearn.model_selection import train_test_split
X_train_iris, X_test_iris, y_train_iris, y_test_iris = train_test_split(iris.data, iris.target, test_size=0.2)
X_train_diab, X_test_diab, y_train_diab, y_test_diab = train_test_split(diab.data, diab.target, test_size=0.2)
```

```
In [17]: import sklearn
import sklearn.linear_model
log_reg=sklearn.linear_model.LogisticRegression(max_iter=200)
log_reg.fit(X_train_iris, y_train_iris)
y_pred=log_reg.predict(X_test_iris)
print ('accuracy: ',log_reg.score(X_test_iris,y_test_iris))
```

accuracy: 0.9666666666666667

```
In [18]: import sklearn.metrics as skm
mat=skm.confusion_matrix(y_pred,y_test_iris)
disp=skm.ConfusionMatrixDisplay(mat)
disp.plot()
|
```

Out[18]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7fe43913adf0>

